

News Headline Classification

Hengxin Wu

Muwen Huang

Chenyue Jin

Jianqiao Zhang



GEORGETOWN UNIVERSITY



Catalog

- Introduction
- **Headline Classification in English News**
 - LSTM
 - CNN
 - BERT
 - HAN
- **Headline Classification in Chinese News**
 - LSTM
- Conclusion

Introduction

The Washington Post



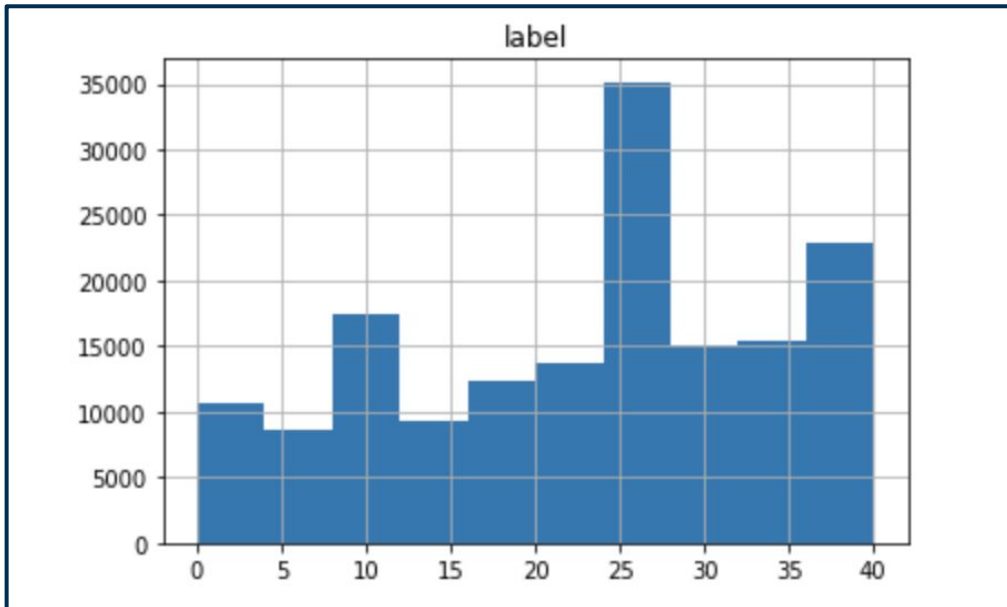
News Category Dataset v2

- **Dimension:** 200853 rows and 40 unique labels
- **Date range:** from 2012-01-28 to 2018-05-26
- **Data fields:** category, headline, authors, link, short description, date
- **Biggest drawback:** imbalanced labels


Data Source: <https://www.kaggle.com/datasets/yazansalameh/news-category-dataset-v2>

Label Distribution - English News with all categories

Train




Test




Headline Length Distribution - English News with all categories

Train




Test




Label Distribution - English News with top 10 categories

Train




Test




Headline Length Distribution - English News with top 10 categories

Train



Test







LSTM

LSTM

- Long Short Term Memory networks are a special kind of RNN, capable of learning long-term dependencies.
- A chain of repeating modules of neural network.
- Have four neural network layers, interacting in a very special way.



LSTM Model

```
model = Sequential()
model.add(layers.Embedding(vocab_size, embedding_dim,
                           weights=[embedding_matrix],
                           input_length=maxlen,
                           #maxlen
                           trainable=True))
model.add(LSTM(64, return_sequences=True))
model.add(LSTM(64, return_sequences=True))
model.add(LSTM(64))
model.add(layers.Dense(40, kernel_regularizer=l2(0.05), activation='softmax'))

opt = Adam(lr=0.005)
model.compile(optimizer=opt,
              loss='categorical_crossentropy',
              metrics=['accuracy'])


history = model.fit(X_train, y_train,
                    epochs=100,
                    validation_data=(X_test, y_test),
                    batch_size=128,
                    class_weight = class_weights)
loss, accuracy = model.evaluate(X_train, y_train, verbose=False)
print("Training Accuracy: {:.4f}".format(accuracy))
loss, accuracy = model.evaluate(X_test, y_test, verbose=False)
print("Testing Accuracy: {:.4f}".format(accuracy))
```

- Keras
- Sequential Layer
- Embedding Layer
- LSTM Layer
- Dense Layer

Train Accuracy VS Test Accuracy


Full

- Accuracy for cross validation is around 0.5



Partial (Top 10 labels)

- Accuracy for cross validation is around 0.7





1D-CNN

1D-CNN

- 2D CNNs use 2D convolutional kernels to predict the segmentation map for a single slice.
- 1D Convolutional Neural Networks are similar to 2D CNN.
- 1D Convolutional Neural Networks are used mainly used on text and 1D signals.




Figure 6.26 How 1D convolution works: each output timestep is obtained from a temporal patch in the input sequence.

1D-CNN Model


```
embedding_dim = 50
model = Sequential()
model.add(layers.Embedding(vocab_size, embedding_dim, input_length=maxlen))
model.add(layers.Conv1D(128, 5, activation='relu'))
model.add(layers.GlobalMaxPooling1D())
model.add(layers.Dense(64, kernel_regularizer=l2(0.005), activation='relu'))
model.add(layers.Dense(40, kernel_regularizer=l2(0.005), activation='softmax'))
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(X_train, y_train,
                   epochs=10,
                   verbose=True,
                   validation_data=(X_test, y_test),
                   batch_size=128,
                   class_weight = class_weights)
loss, accuracy = model.evaluate(X_train, y_train, verbose=False)
print("Training Accuracy: {:.4f}".format(accuracy))
loss, accuracy = model.evaluate(X_test, y_test, verbose=False)
print("Testing Accuracy: {:.4f}".format(accuracy))
#plot_history(history,name='CNN')
```


- Embedding Layer
- Conv1D Layer
- GlobalMaxPooling1D Layer
- Dense Layer

Train Accuracy VS Test Accuracy

Full



Partial






BERT

BERT

- A transformer-based model that uses self-attention as a primary means to express the relationship between variables in a sequence.
- Since Bert is originally designed for generating language models, it utilizes the [CLS] token and the [SEP] token to represent the beginning and the end of each sentence.
- For classification tasks, a classification layer can be added on top of the output for the [CLS] token.



Bert Model

```
for i, batch in enumerate(batches):
    X, Y, _ = batch
    inputs = torch.tensor(X['input_ids'], device=device)
    attmsk = torch.tensor(X['attention_mask'], device=device)
    labels = torch.tensor(Y, device=device)
    batch = {'input_ids': inputs,
            'attention_mask': attmsk,
            'labels': labels}
    with grad_mode():
        outputs = model(**batch)
        embeds.append(outputs[-1][1][:, 0, :].squeeze().detach().cpu()) # only take CLS tokens
        loss = outputs.loss
        if train:
            loss.backward() # Back Propagation
            optimizer.step() # update optimizer (Gardient Descent)
            lr.step() # update learning rate
            optimizer.zero_grad() # clears old gradients from the last step
        logits = outputs.logits
        Yhat = torch.argmax(logits, dim=-1) # perform on last dimension
        preds.append(Yhat)
```

Training loop from scratch

- Retrieve input ids, attention mask, and labels from batch
- Apply pretrained model
- Save CLS embeddings
- Perform gradient descent for training batches
- Get the predictions

Using all categories

Training set

	loss	f1-score	accuracy
epoch 1	1.550	0.412	0.600
epoch 2	1.201	0.514	0.677

Test set

	loss	f1-score	accuracy
epoch 1	1.221	0.506	0.666

Using top 10 categories

Training set

	loss	f1-score	accuracy
epoch 1	0.618	0.765	0.812
epoch 2	0.430	0.834	0.869

Test set

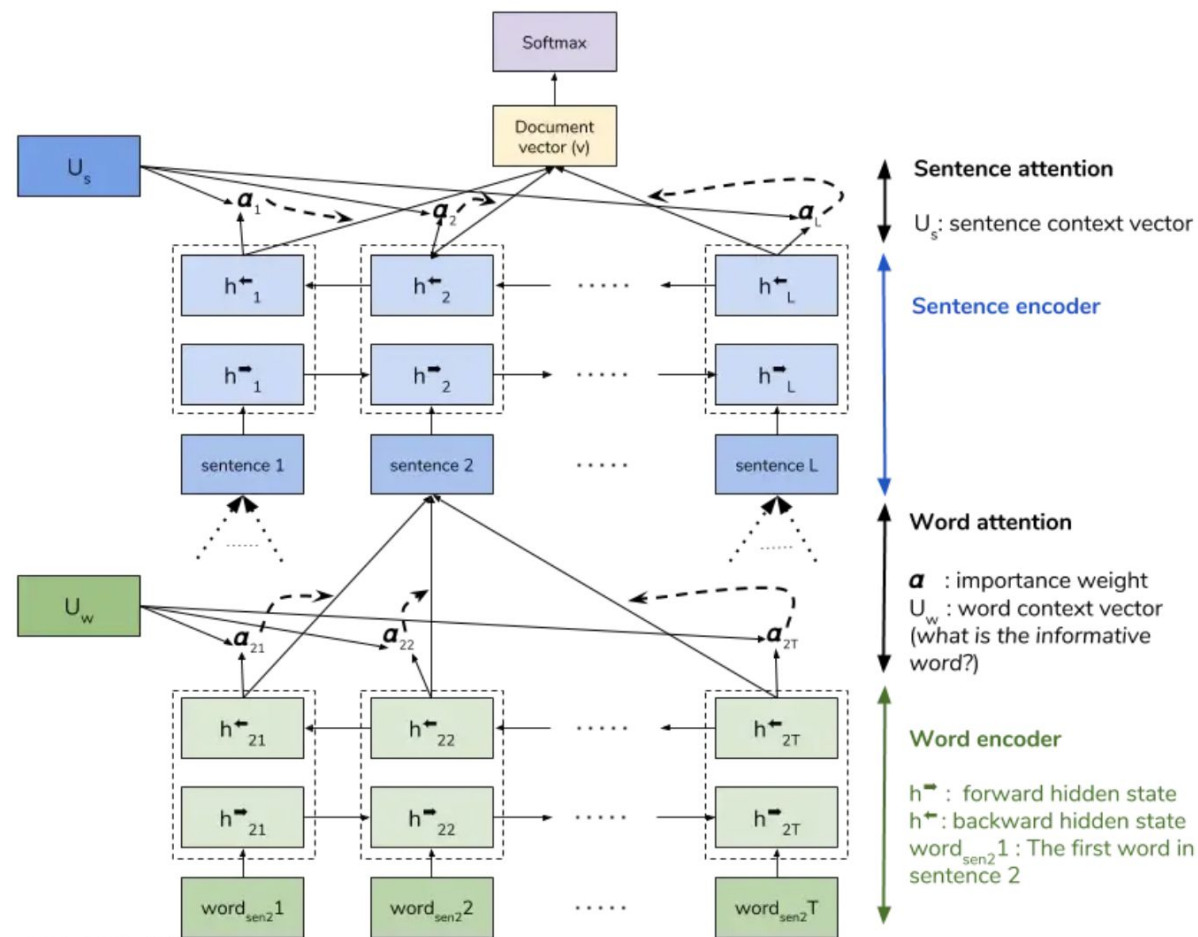
	loss	f1-score	accuracy
epoch 1	0.463	0.817	0.854



Hierarchical Attention Networks

HAN

It uses stacked recurrent neural networks on word level followed by attention model.



bofin.babu@CloudSek

HAN Structure

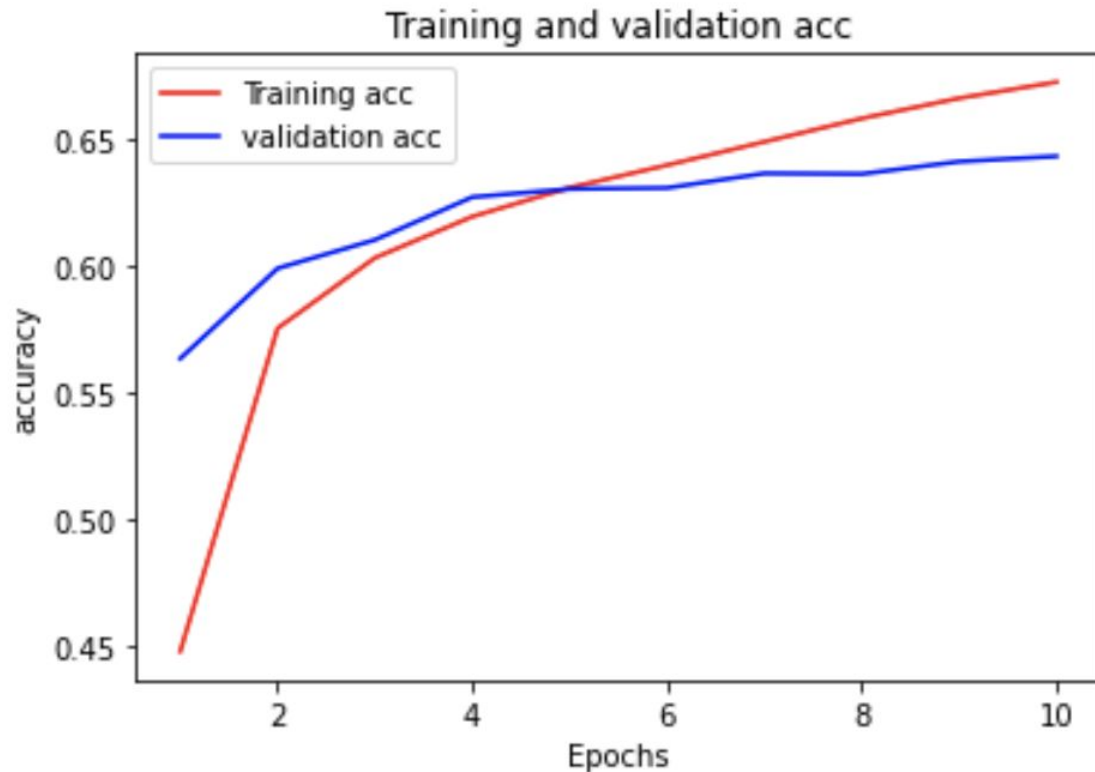
HAN Model

```
l2_reg = l2(0.0000001)
word_input = Input(shape=(max_senten_len,), dtype='float32')
word_sequences = embedding_layer(word_input)
word_lstm = Bidirectional(LSTM(150, return_sequences=True, kernel_regularizer=l2_reg))(word_sequences)
word_dense = TimeDistributed(Dense(200, kernel_regularizer=l2_reg))(word_lstm)
word_att = AttentionWithContext()(word_dense)
wordEncoder = Model(word_input, word_att)


sent_input = Input(shape=(max_senten_num, max_senten_len), dtype='float32')
sent_encoder = TimeDistributed(wordEncoder)(sent_input)
sent_lstm = Bidirectional(LSTM(150, return_sequences=True, kernel_regularizer=l2_reg))(sent_encoder)
sent_dense = TimeDistributed(Dense(200, kernel_regularizer=l2_reg))(sent_lstm)
sent_att = Dropout(0.5)(AttentionWithContext()(sent_dense))
preds = Dense(40, activation='softmax')(sent_att)
model = Model(sent_input, preds)
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=[ 'acc' ])
```


Training Accuracy VS Testing Accuracy

Full



Partial





Headline Classification in Chinese News (LSTM)

Chinese News Headline Classification

- 38000+ news headlines, 15 categories (labels) from a Chinese news application

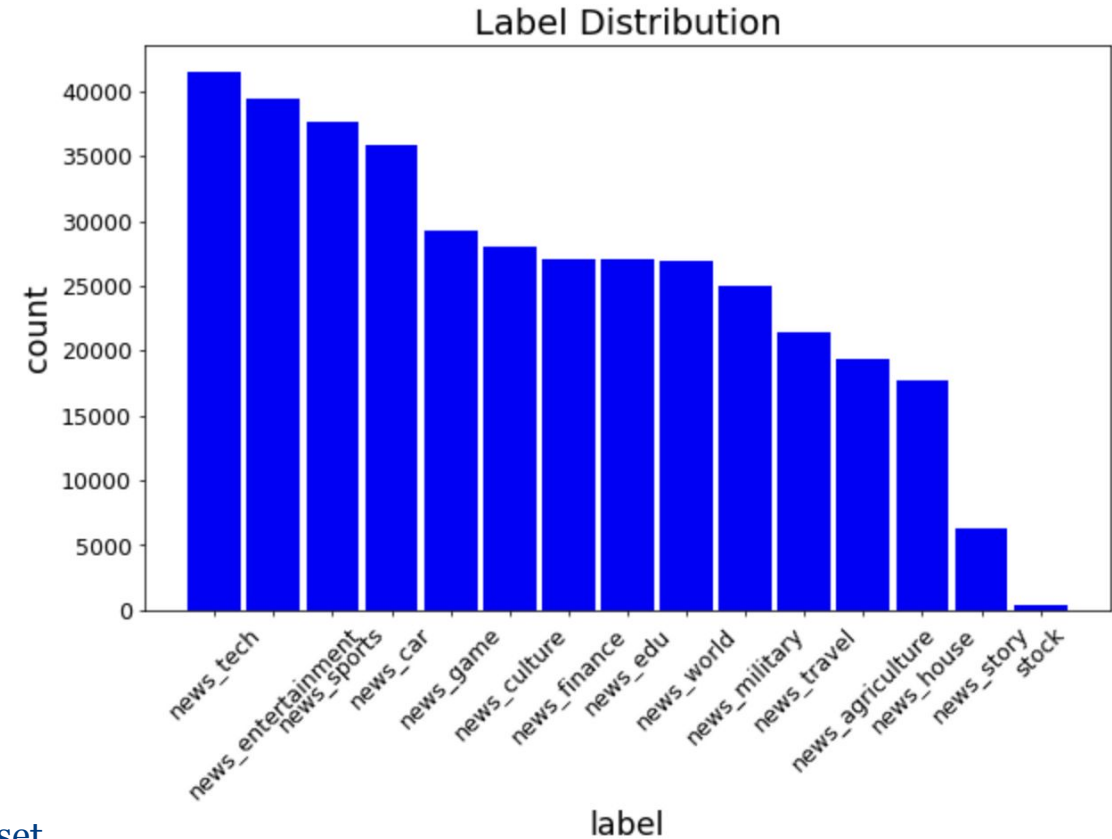


[TouTiao.com](https://www.toutiao.com)

Owner: ByteDance

Data Source:

<https://github.com/aceimnorstuvwxyz/toutiao-text-classfication-dataset>



Chinese News Headline Classification


- Jieba package for Chinese text segmentation
- Full Mode

	label	news
18328	news_house	轻 奢 生活 百花 启幕 碧桂园 桂园 百花 里 产品 发布 发布会 璀璨 绽放 清镇 碧桂...
242156	news_travel	倒计 倒计时 计时 天 襄阳 公里 一级 路标 标准 西三环 三环 环线 即将 开建 牛 首...
177382	news_finance	有没有 没有 些小 小众 质量 品牌
369178	news_entertainment	应 采 怀 时 接地 地气 霍思燕 肚子 更 可能 当年 王菲 小时 时代 怀孕 霍思燕 身...
88378	news_edu	宝宝 念 小学 二年 二年级 年级 成绩 差 爱 学习 办法 提高 学习 兴趣
162753	news_edu	九江 职 院 第八 第八届 八届 全国 国大 大学 大学生 学生 生机 机械 创新 新设 设...
192158	news_sports	克服 习惯 习惯性 惯性 崴 脚
111159	news_finance	希腊 房子 快 买光 传言 究竟 竟是 真是 真是假 中远 集团 希腊 经济 比雷埃夫斯 夫...
338168	news_finance	岁 股市 女 奇才 豪言 研 股市 尾盘 买入 入法 半年 赚 千万 收藏 尾盘 选股 超...
258235	news_finance	龙头 特 停 卡位 龙头 恒瑞 恒瑞医药 医药 新 题材 贵州 复星 医药 新 医药

Example:

孩子念小学 **二年级** 的时候.....

When the child was in the second grade



LSTM Model

```
MAX_NB_WORDS = 50000 # Set the most frequently used 50,000 words
EMBEDDING_DIM = 100 # Set the dimensionality of the Embedding layer


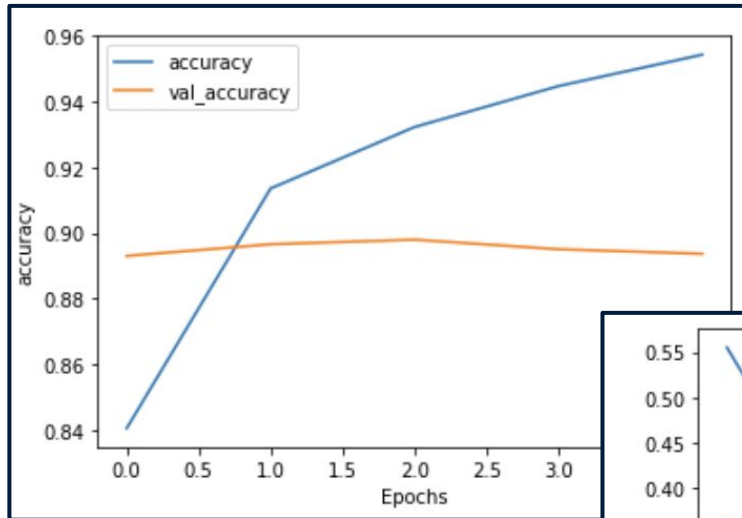
import tensorflow as tf
model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Embedding(MAX_NB_WORDS, EMBEDDING_DIM, input_length=X.shape[1]))
# dropout randomly and independently sets some elements to zero,
# while SpatialDropout1D randomly sets all elements to zero for a particular latitude
model.add(tf.keras.layers.SpatialDropout1D(0.2))
model.add(tf.keras.layers.LSTM(100, dropout=0.2, recurrent_dropout=0.2))
# Output layer contains 15 fully-connected layers for classification,
# and the activation function is set to softmax
model.add(tf.keras.layers.Dense(15, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())

# Start training
from keras.callbacks import EarlyStopping
# Set 5 training cycles
epochs = 5
# Specify the number of samples to be included in each batch during gradient descent
batch_size = 64
# validation_split specifies 10% data in the training set as the validation set
history = model.fit(X_train, Y_train, epochs=epochs, batch_size=batch_size, validation_split=0.1,
                    callbacks=[EarlyStopping(monitor='val_loss', patience=3, min_delta=0.01)])
# callbacks: prevent overfitting and stop training early
model.save(r'./model.h5') # save the model
```

- Keras
- Sequential Layer
- Embedding Layer
- Dropout Layer
- LSTM Layer
- Dense Layer

Train Accuracy VS Test Accuracy

- Accuracy for cross validation is around 0.89 and the loss is around 0.37
- Accuracy for test set: 88.8%



	precision	recall	f1-score	support
news_tech	0.88	0.89	0.88	2759
news_entertainment	0.90	0.91	0.91	3938
news_sports	0.94	0.95	0.95	3765
news_car	0.82	0.82	0.82	2716
news_game	0.91	0.92	0.92	1782
news_culture	0.93	0.93	0.93	3514
news_finance	0.90	0.91	0.90	2615
news_edu	0.88	0.87	0.87	4248
news_world	0.89	0.86	0.87	2445
news_military	0.85	0.83	0.84	2191
news_travel	0.85	0.83	0.84	2706
news_agriculture	0.87	0.89	0.88	1957
news_house	0.91	0.93	0.92	2924
news_story	0.33	0.07	0.12	42
stock	0.80	0.78	0.79	662
accuracy			0.89	38264
macro avg	0.84	0.83	0.83	38264
weighted avg	0.89	0.89	0.89	38264



Conclusion

CONCLUSION

- Bert has significantly higher test accuracies among all the models for both all categories and the top 10 categories.
- HAN also has improved accuracies than LSTM and 1D-CNN.
- Note that BERT and HAN can take a lot more time to train!
- Why is the accuracy rate in Chinese higher?
 - Less categories/labels (40 compared to 15).
 - Chinese news headlines have more tokens due to Chinese text segmentation.



Thanks for watching!